

Eigen Labs

EigenDA Database (LittDB) Security Review

:Final Report

May 17th, 2025

Revision 1.1

Prepared for *Eigen Labs*

Prepared by [ChainLight](#) of Theori

Theori, Inc. (“We”) is acting solely for the client and is not responsible to any other party. Deliverables are valid for and should be used solely in connection with the purpose for which they were prepared as set out in our engagement agreement. You should not refer to or use our name or advice for any other purpose. The information (where appropriate) has not been verified. No representation or warranty is given as to accuracy, completeness or correctness of information in the Deliverables, any document, or any other information made available. Deliverables are for the internal use of the client and may not be used or relied upon by any person or entity other than the client. Deliverables are confidential and are not to be provided, without our authorization (preferably written), to entities or representatives of entities (including employees) that are not the client, including affiliates or representatives of affiliates of the client.

Table of Contents

1. Executive Summary	3
2. Overview	4
2.2. Project Scope	4
2.3. Revision History	4
2.4. Threat Actors and Scenarios	4
2.5. Testing Methodology	5
2.6. Findings Breakdown by Severity	6
3. Findings	7
#1 HashKey needs a stronger permutation	8
#2 Use unsafe.String	10
#3 Potential Goroutine and Channel Leak from Early Return	12
#4 Race Condition in Flush Order Between Shard and Key File Control Loops	15
#5 Under-Provisioned Slice Capacity for batchData Leading to Performance Degradation	18
#6 Validator Crash Due to Mismatched Bundles and Metadata Size from Malicious Relay	20
#7 Shadowed Loop Variable Affecting Code Readability	22
#8 Race Condition in storeBatchLittDB May Violate Validation Invariants	23
#9 Handling Empty or Corrupt Segments from Repeated Restarts	26
4. Conclusion	28

1. Executive Summary

From April 28, 2025, to May 13, 2025, Theori executed a security review of LittDB, a highly specialized embedded key-value store designed for EigenDA and optimized for high write throughput, low read latency, low memory usage, write-once/never-update data, and data deletion exclusively via a Time-to-Live (TTL) mechanism. The review concentrated on LittDB's implementation and its integration with EigenDA's validator/operator, while excluding external components and configurations. Comprehensive scope details, including pertinent pull request URLs, are provided within the "Project Overview" section. This review was conducted by two security engineers over a period of two and a half weeks.

The goal of this security review is to verify that LittDB operates correctly and securely under various conditions. This includes handling abnormal data processing, write durability, crash safety, and prevention of deadlocks and race conditions. The threat actors and scenarios considered are malicious senders, distributors, relays, and validators, as well as natural disasters. File system corruption and CPU bugs were excluded. Due to the protocol's nature, external attacker inputs are extremely limited, so the primary focus was on system resilience during errors caused from natural disasters¹. The review employed a testing methodology that included both manual code review as the primary method and fault injection testing. Fault injection testing utilized Go tests, Cgroups memory limits, and simulated sudden termination due to SIGKILL signals.

Several findings were identified during the review. A critical finding ([#1](#)) was that the current linear permutation in hashing algorithm allows attackers to manipulate shard mappings. Replacing this with a stronger, non-linear permutation function like SipHash is recommended. A high-severity race condition ([#4](#)) between the shard control loop and key file control loop was also identified, which could result in data inconsistency upon system restart. A medium-severity validator crash ([#6](#)) was identified where a malicious relay could send a response with mismatched bundles and metadata, causing an out-of-bounds array access. This has been addressed by implementing a length check before processing the response.

We recommend addressing the identified findings based on their severity to ensure the security and robustness of LittDB within the EigenDA ecosystem.

The review focused on validating whether the system operates correctly and safely under a variety of conditions. For example, we ensured that the system remains robust when handling malformed chunks, preventing such inputs from compromising overall stability. More broadly, the review assessed whether key assumptions related to correctness and safety hold true, such as guarantees around write durability, crash resilience, and the avoidance of deadlocks and race conditions.

¹ This includes system crashes that can result from unforeseen events, such as data center disasters.

2. Overview

Security review of LittDB and its EigenDA integration (validator/server) was performed by two security engineers over 2.5 weeks (April 28 - May 13, 2025), as detailed in provided pull requests. The scope included analyzing threat actors and scenarios, with a focus on write durability, crash resistance, and preventing malicious disruption of LittDB. Considered threats included natural disasters, while disk file corruptions and CPU bugs were excluded. Testing involved manual code review and fault injection using Go tests, Cgroups memory limits, and the sudden termination due to SIGKILL.

2.2. Project Scope

The primary focus was directed towards the following Pull Requests pertaining to the LittDB implementation and its integration with EigenDA:

- LittDB: <https://github.com/Layr-Labs/eigenda/pull/1280>
- LittDB Integration: <https://github.com/Layr-Labs/eigenda/pull/1429>

Additionally, patch Pull Requests addressing specific findings were also taken into account.

2.3. Revision History

- May 13, 2025 Initial Report
- May 17, 2025 Reflect the team's response to issue #4

2.4. Threat Actors and Scenarios

The security analysis considered potential threats posed by various actors interacting with the LittDB implementation and its integration within the EigenDA ecosystem. These include:

- **Malicious Sender:** An unidentified entity capable of transmitting arbitrary data blobs to the EigenDA pipeline.
- **Malicious Disperser / Relay:** A compromised disperser with the ability to manipulate signed data, and a compromised relay that can disseminate crafted data, potentially leading to denial-of-service conditions before validation.
- **Malicious Validators:** Validator (Operator) nodes that process signed payloads from dispersers and may subsequently propagate or replay these payloads to other validators within the network.
- **Natural Disasters:** Unforeseen environmental events with the potential to cause sudden and irreparable system failures, such as solar flares inducing hardware anomalies². It is important to note that filesystem-level corruption is outside the scope of this review, as

² CPU bugs ([ref](#)), ECC RAM crashes, and other issues are considered outside the scope of this review.

the underlying kernel is assumed to ensure file system integrity³ upon successful completion of flush operations.

- a. The primary focus regarding natural disasters is ensuring crash safety, specifically the crash durability of flushed data.

Further consideration was given to advanced attack vectors, such as hot shard attacks and attempts to circumvent cryptographic salt mechanisms. Additionally, the analysis encompassed potential issues arising from unintentional process termination and other operational scenarios that could compromise system stability and data integrity.

2.5. Testing Methodology

The primary goal of this security review and testing process was to ensure a high level of confidence in the core security and reliability features of the EigenDA LittDB system. In particular, our focus was on carefully analyzing and validating the following key properties:

- **Write Durability:** Data acknowledged after a FLUSH operation must persist despite unexpected power loss, with recovery taking constant time $O(1)$.
- **Crash Safety:** Following any crash, the database must restart in a recoverable state without replay causing mutations or hangs.
- **Deadlock & Race Freedom:** The system must be free of hard locks (e.g., channel locks) and goroutine races between the key-file writer and the shard writer.

The security review process encompassed the following activities:

- **Threat Modeling:** Identifying potential actors and adversarial scenarios, including power loss, storage tampering, and byzantine replicas, and mapping corresponding security controls.
- **Manual Code Review:** Performing a detailed, line-by-line examination of the codebase for security vulnerabilities.
- **Adversarial Fault Injection:** Employing a manual chaos engineering approach to introduce targeted syscalls, disk faults, and network partitions.

The primary testing methodology employed was manual code review, supplemented by fault injection testing. Fault injection was carried out using Go-based tests, enforced Cgroups memory limits, and simulated abrupt termination through SIGKILL signals.

As part of our threat modeling and testing efforts, we evaluated a variety of potential risks to the security and reliability of EigenDA's LittDB system. This included scenarios such as natural disasters, which could cause power loss or physical hardware damage. However, it is important to note that this review explicitly excluded vulnerabilities related to underlying disk file corruption not caused by system crashes, as well as fundamental CPU-level bugs. These areas were considered outside the scope of the current engagement.

³ This review assumes that even with crashes, the file system remains intact.

2.6. Findings Breakdown by Severity

This section details the security review findings, categorized by their severity level.

Category	Count	Findings
High	2	<ul style="list-style-type: none">• Finding #1• Finding #4
Medium	1	<ul style="list-style-type: none">• Finding #6
Low	2	<ul style="list-style-type: none">• Finding #3• Finding #9
Info	4	<ul style="list-style-type: none">• Finding #2• Finding #5• Finding #7• Finding #8

3. Findings

This table summarizes the security audit findings. Detailed information for each finding can be found below.

#	Title	Severity	Status
1	HashKey needs a stronger permutation	High	Fixed
2	Use unsafe.String	Info	Fixed
3	Potential Goroutine and Channel Leak from Early Return	Low	Fixed
4	Race Condition in Flush Order Between Shard and Key File Control Loops	High	WIP
5	Under-Provisioned Slice Capacity for batchData Leading to Performance Degradation	Info	Fixed
6	Validator Crash Due to Mismatched Bundles and Metadata Size from Malicious Relay	Medium	Fixed
7	Shadowed Loop Variable Affecting Code Readability	Info	Fixed
8	Race Condition in storeBatchLittDB May Violate Validation Invariants	Medium	Won't Fix
9	Handling Empty or Corrupt Segments from Repeated Restarts	Info	WIP

#1 HashKey needs a stronger permutation

Summary

The current linear permutation in *util.HashKey* allows attackers to map to the same shard or lock index, even without knowing the salt. The XOR operation does not provide enough diffusion, and the permhash is easily reversible. To mitigate this, stronger permutation alternatives like SipHash are being evaluated to balance compute cost and collision resistance. Additionally, hot-shard monitoring is considered for long-term detection and rebalancing of busy shards.

Description

The vulnerability lies in the *util.HashKey* function, where the current permutation is linear. This linearity means that an attacker can predict the shard or lock index, even without knowledge of the salt. The existing XOR-based permutation lacks sufficient diffusion and is reversible, making it possible to manipulate the mapping.

```
Go
//
// https://github.com/Layr-Labs/eigenda/blob/81c86cb0509a9fc8abc371748b84c953b0b77
// 80e/litt/util/hashing.go#L18C1-L59C2
func Perm64(x uint64) uint64 {
    // This is necessary so that 0 does not hash to 0.
    // As a side effect this constant will hash to 0.
    x ^= 0x5e8a016a5eb99c18

    x += x << 30
    x ^= x >> 27
    x += x << 16
    x ^= x >> 20
    x += x << 5
    x ^= x >> 18
    x += x << 10
    x ^= x >> 24
    x += x << 30
    return x
}

// Perm64Bytes hashes a byte slice using perm64.
func Perm64Bytes(b []byte) uint64 {
    x := uint64(0)

    for i := 0; i < len(b); i += 8 {
        var next uint64
        if i+8 <= len(b) {
            // grab the next 8 bytes

```



```

        next = binary.BigEndian.Uint64(b[i:])
    } else {
        // insufficient bytes, pad with zeros
        nextBytes := make([]byte, 8)
        copy(nextBytes, b[i:])
        next = binary.BigEndian.Uint64(nextBytes)
    }
    x = Perm64(next ^ x)
}

return x
}

// HashKey hashes a key using perm64 and a salt.
func HashKey(key []byte, salt uint32) uint32 {
    return uint32(Perm64(Perm64Bytes(key) ^ uint64(salt)))
}

```

Impact

High

This vulnerability can lead to uneven distribution of data or processing load across shards or lock indices. Attackers could exploit this vulnerability to target specific shards, potentially leading to a denial-of-service or rendering the sharding factor ineffective..

Recommendation

The primary recommendation is to replace the current linear permutation with a stronger, non-linear permutation function. SipHash is suggested as a potential alternative, but the solution should be evaluated to find a balance between computational cost and collision resistance. In addition, consider implementing hot-shard monitoring to detect and rebalance shard loads.

Remediation

The issue was addressed by integrating the suggested SipHash algorithm.

PR: <https://github.com/Layr-Labs/eigenda/pull/1506>

#2 Use unsafe.String

Summary

Go 1.21 allows for replacing unsafe string conversions with *unsafe.String*, resulting in cleaner and safer code. Data corruption is a very unlikely outcome of this bug, occurring only if the compiler's escape analysis fails.

Description

The current code employs unsafe string conversions. Go 1.21 introduces *unsafe.String*, a potentially cleaner and safer alternative. While current compiler escape analysis makes data corruption unlikely, it's not impossible. A compiler bug or future code modifications could still introduce vulnerabilities leading to data corruption.

```
Go
func UnsafeBytesToString(b []byte) string {
    return *(*string)(unsafe.Pointer(&b))
}
```

Impact

Info

Refactoring with *unsafe.String* improves code readability and safety. While the risk of compiler escape analysis causing data changes is low, it should be addressed to avoid potential issues.

Recommendation

Consider replacing it with the *unsafe.String*.

Remediation

The issue was addressed by replacing it with *unsafe.String*.

```
Go
//
// https://github.com/Layr-Labs/eigenda/blob/63e0e5e063e0d8f21aa42d12f5b5780510b37
// 1dd/litt/util/unsafe_string.go#L7
func UnsafeBytesToString(b []byte) string {
```

```
    if len(b) == 0 {  
        return ""  
    }  
    return unsafe.String(&b[0], len(b))  
}
```

PR: <https://github.com/Layr-Labs/eigenda/pull/1507>

#3 Potential Goroutine and Channel Leak from Early Return

Summary

A potential goroutine leak exists within a for-loop due to an early return that might occur before all data from the *writeCompleteChan* output channel is processed. If the channel buffer fills, the corresponding goroutine will become permanently blocked, leading to unbounded memory consumption as unreferenced goroutines are not garbage collected (A). Additionally, a channel leak can occur because more channel outputs are sent than expected when the Put operation returns an error (B).

Description

The issue occurs when an early return statement is executed inside a for loop that is designed to process data from an output channel, *writeCompleteChan*. If the loop terminates prematurely (A), before all expected data has been read from this channel, the goroutine responsible for sending data to *writeCompleteChan* might remain blocked. This blockage happens if the channel's buffer fills up, as the goroutine will wait indefinitely for a receiver that no longer exists due to the early return.

```
Go
//
https://github.com/Layr-Labs/eigenda/blob/3e918b09df7962ff772b0cedf27dc508603e189e/node/validator\_store.go#L370C1-L409C3
    writeCompleteChan := make(chan error, len(batchData))
    for _, batchDatum := range batchData {
        bundleKeyBytes := batchDatum.BundleKey
        bundleData := batchDatum.BundleBytes

        go func() {
            // Grab a lock on the hash of the blob. This protects
            against duplicate writes of the same blob.
            lockIndex := uint64(util.HashKey(bundleKeyBytes[:],
s.duplicateRequestSalt))
            s.duplicateRequestLock.Lock(lockIndex)
            defer s.duplicateRequestLock.Unlock(lockIndex)

            exists, err := s.chunkTable.Exists(bundleKeyBytes[:])
            if err != nil {
                writeCompleteChan <- fmt.Errorf("failed to check
existence: %v", err)
                return
            }

            if exists {
                // Data is already present, no need to write it
                again.
```

```

        writeCompleteChan <- nil
        return
    }

    err = s.chunkTable.Put(bundleKeyBytes, bundleData)
    if err != nil {
        writeCompleteChan <- fmt.Errorf("failed to put data:
%v", err) // HERE (B)
    }

    writeCompleteChan <- nil
}()

size += uint64(len(bundleKeyBytes) + len(bundleData))
}

for i := 0; i < len(batchData); i++ {
    err := <-writeCompleteChan
    if err != nil {
        return 0, err // HERE (A)
    }
}

```

Goroutines that are blocked and have no remaining references in the program cannot be reclaimed by the Go runtime's garbage collector. Consequently, the memory allocated to such goroutines and their associated resources, including the channel itself, will not be freed. Over time, if this scenario repeats, it can lead to a continuous increase in memory consumption, potentially exhausting system resources.

Impact

Low

Although unlikely under typical usage, a potential risk exists for unbounded memory growth. Early returns could cause channel buffer saturation, leading to an accumulation of blocked, unreferenced goroutines. Over time, this memory leak could degrade system performance and stability, potentially resulting in an out-of-memory error and application crash. The probability of this issue occurring is considered very low.

Recommendation

To address this potential leak, two primary actions are recommended:

1. Prevent Channel Buffer Saturation: Introduce a return statement at the point where the early exit occurs within the for-loop. This ensures that if the loop is exited prematurely,

the function also returns, preventing further operations that might lead to the *writeCompleteChan* buffer filling up from the perspective of the current execution path.

2. Drain Remaining Channel Inputs: Implement a separate goroutine to drain any remaining inputs from the *writeCompleteChan*. This can be achieved by launching a new goroutine specifically tasked with reading and discarding the outstanding items from the channel. For instance:

```
Go
go func(remaining) {
    for i:=0; i<remaining; i+=1 {
        _ := <- writeCompleteChan
    }
}(len(batchData) - i))
```

3. This ensures that the sending goroutine can complete its operations and will not remain blocked, allowing it to be properly garbage-collected. The argument *len(batchData) - i* would represent the number of items yet to be read from the channel at the point of the early return.

Remediation

The return statement has been implemented as recommended. Additionally, the channel draining logic was modified to guarantee that all items are processed, achieving the same outcome as our suggestion. This ensures that even if the channel's capacity is insufficient for all data, insertions will not be blocked.

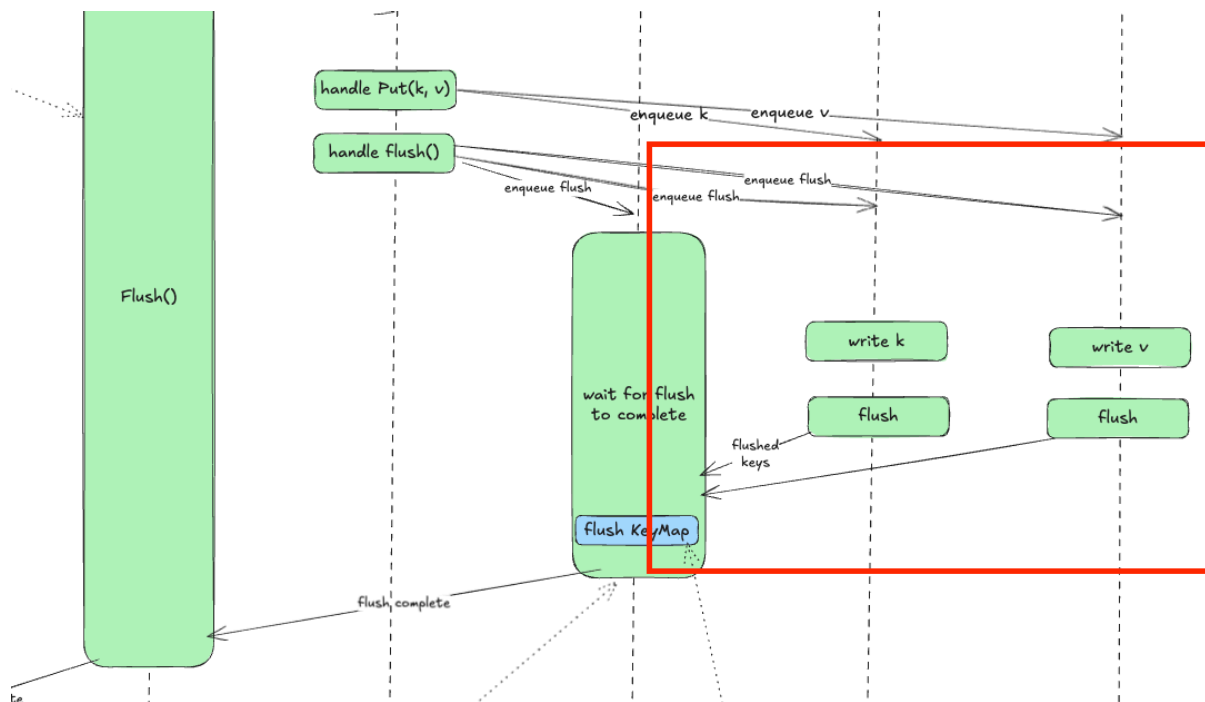
PR: <https://github.com/Layr-Labs/eigenda/pull/1535>

#4 Race Condition in Flush Order Between Shard and Key File Control Loops

Summary

A race condition exists between the *shardControlLoop* and *keyFileControlLoop* due to their independent invocation of the *Flush()* operation. If a system crash occurs after the key file is persisted but before all associated shard files are flushed, the *reloadKeymap* function may incorrectly validate keys. This happens because *reloadKeymap* assumes that once a valid key-shard pair is found, all subsequent entries are also valid, potentially leading to "dangling" keys that appear to exist but lack their corresponding shard data, thereby compromising data consistency.

Description



The executions/operations within the red box are independent of each other.

The *shardControlLoop* and *keyFileControlLoop* are responsible for managing shard data and key file persistence, respectively. Each loop calls the *Flush()* function independently to write its data to disk. A critical race condition can arise if the system crashes at a specific point: after the key file (managed by *keyFileControlLoop*) has been successfully flushed, but before all shard files (managed by *shardControlLoop*) related to the keys in that key file have been flushed. This scenario can occur, for example, if the last key's shard file was flushed, but other preceding shard files were not.

Upon system restart, the *reloadKeymap* function is executed to reconstruct the in-memory key map. This function operates under the assumption that "once we find a valid pair, all subsequent entries are assumed good." If *reloadKeymap* encounters keys in the persisted key file for which the corresponding shard data was not flushed due to the crash, it may still mark these keys as valid. This results in "dangling" keys within the keymap: these keys will pass existence checks (e.g., an *Exists* check will return true), but their underlying shard data is missing or incomplete. Consequently, attempts to use these keys, or to re-add a key with the same identifier, may fail or lead to inconsistent states.

Impact

High

The primary impact of this race condition is data inconsistency. The presence of "dangling" keys means that the system believes certain data exists and is accessible when, in fact, its constituent shard data is not fully persisted or is missing. This can lead to several issues:

- Validators may fail or exhibit unexpected behavior when trying to access data linked to a dangling key.
- Writing new data with a dangling key might fail because the system's *Exists* check could incorrectly identify the key as occupied.
- The overall integrity of the data is undermined.

Recommendation

To mitigate this issue, the following measures are recommended:

1. Enforce Strict Flush Ordering: Implement a mechanism to ensure that shard data is always flushed to disk before the corresponding key file is flushed. This reduces the window where a crash could leave the key file referencing unpersisted shards.
 - Caveat: It is acknowledged that this may not be a 100% foolproof solution, as the operating system or Go runtime may automatically flush buffers under certain circumstances (e.g., if a buffer fills up), potentially out of the application's direct control.
2. Implement Validation on Reload: Enhance the *reloadKeymap* function to perform a thorough validation for each key. Before marking a key as present in the in-memory key map, the function should verify that the corresponding shard file(s) exist and are intact.
 - If this validation fails for a key, the system should trigger a corrective action, such as rebuilding the specific key entry (if possible from other sources), rolling back the entry from the key map, or logging a critical error for manual intervention.

Remediation

EigenDA Team Response:

I think an argument can be made that this is a zero-probability event, assuming there are no bugs in LevelDB that require us to attempt manual data recovery (in which case, all bets are off since we are contemplating some undefined bug that corrupts arbitrary data).

In order for the race condition to cause problems, we must rebuild the keymap. The keymap is rebuilt in two circumstances:

- 1) the user decides to switch the keymap implementation used for an existing database
- 2) the database is configured to use an in-memory keymap implementation that needs to be rebuilt at boot time

Neither of these conditions are possible unless a validator's operator makes code changes and runs with a modified validator implementation. For the time being, validators will only run with the LevelDB keymap implementation, and we will never attempt a keymap migration until the bug is fixed.

When the bug is eventually fixed, key files that were corrupted prior to the bug fix could still cause keymap migration issues. But this has an easy solution: just wait until all pre-migration data has exceeded its TTL and is deleted (this is ~2 weeks for our use case).

ChainLight Response:

The team is aware of the zero-probability event and has a "TODO" in the code as a temporary measure. They have a plan and strategy in place to fix it, and the upcoming patch will also address finding #9.

#5 Under-Provisioned Slice Capacity for batchData Leading to Performance Degradation

Summary

The *batchData* slice is initially allocated with a capacity equal to *len(rawBundles)*. However, when operating with multiple operator quorums, each *rawBundle* can expand into several *node.BundleToStore* entries. This expansion can cause the actual number of entries to exceed the initial capacity, forcing Go to reallocate and resize the underlying slice. These reallocations can degrade performance.

Description

The *batchData* slice, designed for *node.BundleToStore* elements, is initialized with a capacity equal to the number of *rawBundles*. This can be insufficient when multiple operator quorums are active, as a single *rawBundle* might generate multiple *node.BundleToStore* entries. If the total number of these entries exceeds the initial capacity, Go's automatic reallocation and copying of the slice become computationally expensive.

```
Go
//
https://github.com/Layr-Labs/eigenda/blob/1b1b48777cf9da4a00ab0544584ff09060a03922/node/node\_v2.go#L159-L165
    for i, bundle := range resp.bundles {
        metadata := resp.metadata[i]

        blobShards[metadata.blobShardIndex].Bundles[metadata.quorum], err =
        new(core.Bundle).Deserialize(bundle)
        if err != nil {
            return nil, nil, fmt.Errorf("failed to deserialize
        bundle: %v", err)
        }

        rawBundles[metadata.blobShardIndex].Bundles[metadata.quorum] = bundle // HERE

    // ...

//
https://github.com/Layr-Labs/eigenda/blob/1b1b48777cf9da4a00ab0544584ff09060a03922/node/grpc/server\_v2.go#L179-L189
func (s *ServerV2) validateAndStoreChunks(
    ctx context.Context,
    batch *corev2.Batch,
    blobShards []*corev2.BlobShard,
    rawBundles []*node.RawBundles,
    operatorState *core.OperatorState,
    batchHeaderHash [32]byte,
```

```
        probe *common.SequenceProbe,  
    ) error {  
  
        batchData := make([]*node.BundleToStore, 0, len(rawBundles)) // HERE
```

Impact

Info

Insufficiently sized *batchData* slice capacity primarily causes performance issues. Frequent reallocations due to exceeding capacity during high throughput or with many expanding *rawBundles* lead to performance degradation from memory operations.

Recommendation

To optimize memory usage, pre-calculate the total count of *BundleToStore* elements by iterating through all *rb* in *rawBundles* and summing the lengths of their respective *rb.Bundles*. Subsequently, initialize the *BundleToStore* slice with a capacity equal to this pre-calculated total using *make([]*node.BundleToStore, 0, total)*. This approach avoids redundant memory reallocations during the appending process.

Remediation

The changes implement the recommended capacity calculation, thereby optimizing *batchData* slice allocations.

PR: <https://github.com/Layr-Labs/eigenda/pull/1543>

#6 Validator Crash Due to Mismatched Bundles and Metadata Size from Malicious Relay

Summary

A malicious relay can crash a validator node by sending a response containing more bundles than metadata entries during the download process. This mismatch causes an out-of-bounds array access while the validator attempts to pair bundles with their corresponding metadata, resulting in a panic and subsequent crash.

Description

The bug exists in the code that processes responses from relays. When a validator receives a response, it expects a list of bundles and a corresponding list of metadata entries, presumably of the same length. The code iterates through the received bundles using a loop similar to:

Go

```
for i, bundle := range resp.bundles {  
    metadata := resp.metadata[i] // Out of bound access
```

If a malicious relay crafts a response such that $\text{len}(\text{resp.bundles})$ is greater than $\text{len}(\text{resp.metadata})$, then during the iteration, when the index i becomes equal to or greater than $\text{len}(\text{resp.metadata})$, the access $\text{resp.metadata}[i]$ will result in an index out-of-bounds panic. This situation can occur if a relay intentionally returns more bundle data than the validator requested or than what is described by the *metadata*. The panic is unhandled and will cause the validator process to terminate unexpectedly.

Impact

Medium

A vulnerability exists that could cause a denial-of-service (DoS) on validator nodes, preventing them from participating in the network and performing essential duties like attesting to data availability. This could negatively impact the DA network's health and liveness, particularly if multiple validators or a critical validator are affected. However, because relay servers must be compromised to exploit this vulnerability, the impact is considered medium despite the potential for validator crashes.

Recommendation

To prevent this crash, a check should be implemented before iterating through the *resp.bundles*. The lengths of *resp.bundles* and *resp.metadata* must be asserted to be equal. If the lengths

differ, the function should return an error immediately, rather than proceeding with the potentially unsafe iteration.

```
Go  
len(resp.bundles) == len(resp.metadata)
```

Remediation

The changes implement the recommended length check to prevent the panic.

PR: <https://github.com/Layr-Labs/eigenda/pull/1545>

#7 Shadowed Loop Variable Affecting Code Readability

Summary

Reusing the loop iteration variable *i* in a nested loop, while technically valid in Go due to its scoping rules, negatively impacts code readability and complicates logical tracing for developers. Consider using distinct variable names for clarity in nested loops.

Description

The practice of using the same variable name, *i* in this case, as an iterator in both an outer and an inner loop is a common programming pattern. While Go's block scoping rules mean the inner *i* is a distinct variable, preventing execution errors, this reuse of the same name in nested loops can create confusion for readers of the code. Differentiating between the outer and inner loop counters requires additional cognitive effort.

Impact

Info

The primary impact of this issue is on code readability and maintainability. It does not introduce a functional bug or a security vulnerability.

Recommendation

For code clarity and to prevent confusion, use different names for loop variables in nested loops (e.g., *j*, *idx*, or a descriptive name).

Remediation

The changes rename the shadowed loop variables to enhance code clarity.

PR: <https://github.com/Layr-Labs/eigenda/pull/1544>

#8 Race Condition in *storeBatchLittDB* May Violate Validation Invariants

Summary

LittDB is designed to enforce a one-to-one mapping between keys and values. However, concurrent calls to *chunkTable.Put(key, value)* for the same key but potentially different values within the *storeBatchLittDB* function could result in a nondeterministic outcome, where the final value stored for the key is not guaranteed. The practical impact of this is questioned, given EigenDA's specific key derivation (from KZG commitment of the payload) and an earlier validation stage that should reject key/value pairs where the value's commitment does not match the key's embedded commitment.

Description

storeBatchLittDB has a race condition during concurrent Put operations on the same key. Specifically, if *storeBatchLittDB* executes *chunkTable.Put(A, X)* and *chunkTable.Put(A, Y)* concurrently, the final value associated with key A becomes unpredictable due to the timing of operations, despite a lock intended to prevent simultaneous writes to the same key.

```
Go
func (s *validatorStore) storeBatchLittDB(batchData []*BundleToStore) (uint64,
error) {
    var size uint64

    writeCompleteChan := make(chan error, len(batchData))
    for _, batchDatum := range batchData {
        bundleKeyBytes := batchDatum.BundleKey
        bundleData := batchDatum.BundleBytes

        go func() { // go function's execution sequences are
nondeterministic.
            // Grab a lock on the hash of the blob. This protects
            against duplicate writes of the same blob.
            lockIndex := uint64(util.HashKey(bundleKeyBytes[:],
s.duplicateRequestSalt))
            s.duplicateRequestLock.Lock(lockIndex)
            defer s.duplicateRequestLock.Unlock(lockIndex)

            exists, err := s.chunkTable.Exists(bundleKeyBytes[:])
            if err != nil {
                writeCompleteChan <- fmt.Errorf("failed to check
existence: %v", err)
            }
            return
        }
    }
}
```

```

        if exists {
            // Data is already present, no need to write it
            writeCompleteChan <- nil
            return
        }

        err = s.chunkTable.Put(bundleKeyBytes, bundleData)
        if err != nil {
            writeCompleteChan <- fmt.Errorf("failed to put data:
again.
%v", err)
            return
        }

        writeCompleteChan <- nil
    }()

    size += uint64(len(bundleKeyBytes) + len(bundleData))
}

```

However, an earlier validation step aims to prevent inconsistent data from reaching LittDB. This validation checks if the commitment of a value matches the commitment in its key. The key question is whether this EigenDA-specific pre-validation and key derivation logic (through KZG commitment) effectively prevents conflicting writes (where different values X and Y map to the same key context). If such conflicts are impossible due to upstream logic, the nondeterministic behavior in *storeBatchLittDB* under concurrent Put operations on the same key might not be a practical concern.

Impact

Info

In EigenDA, keys are derived from the KZG commitment of their blobs, which effectively prevents scenarios like using the same key for different values. As a result, while the race condition identified in *storeBatchLittDB* is theoretically valid, it is unlikely to be triggered in EigenDA's LittDB use case. Nevertheless, the issue merits inclusion in the report for reference.

Recommendation

Consider implementing a key existence check immediately before calling `storeBatchLittDB` as an additional safeguard.

Remediation

The analysis indicates that the race condition involving conflicting values for the same key is practically untriggerable under EigenDA's LittDB usage patterns:

- Case 1: Identical key-value pairs. Concurrent calls such as *chunkTable.Put(A, X)* and *chunkTable.Put(A, X)* do not compromise data integrity, as the final state remains consistently (A, X) regardless of execution order.
- Case 2: Identical keys with differing values. A scenario like *chunkTable.Put(A, X)* and *chunkTable.Put(A, Y)* (where $X \neq Y$) could lead to a race condition. However, this is considered highly improbable due to two safeguards:
 - Key derivation from value: Keys are generated using a KZG commitment of the value, meaning different values should naturally produce different keys.
 - Validation logic: The system includes checks to ensure that the value matches the derived key (i.e., the commitment aligns with the data). Any mismatch should be detected and rejected before reaching storage.

Therefore, this race condition would only surface if either a cryptographic hash collision occurs (which is negligible in probability) or if there is a fault in upstream validation. As such, no remediation was implemented for the concurrency issue in *chunkTable.Put* within *storeBatchLittDB*, given the mitigating controls already in place.

#9 Handling Empty or Corrupt Segments from Repeated Restarts

Summary

Fault injection testing revealed that if a validator node encounters a bug causing indefinite restarts, it can repeatedly create zero-byte or effectively empty/corrupt unsealed segments. This behavior raises concerns about inode exhaustion and metadata bloat. A multi-layered defense is recommended, either by reusing the zero-byte segment or avoiding its regeneration altogether.

Description

A bug that triggers infinite restarts causes LittDB to generate numerous 0-byte segments, as the last segment is sealed on each startup.

```
-rw-r--r-- 1 vscode vscode 21 May 13 06:38 6.metadata
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 7-0.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 7-1.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 7-2.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 7-3.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 7.keys
-rw-r--r-- 1 vscode vscode 21 May 13 06:38 7.metadata
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 8-0.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 8-1.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 8-2.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 8-3.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 8.keys
-rw-r--r-- 1 vscode vscode 21 May 13 06:38 8.metadata
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 9-0.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 9-1.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 9-2.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 9-3.values
-rw-r--r-- 1 vscode vscode 0 May 13 06:38 9.keys
-rw-r--r-- 1 vscode vscode 21 May 13 06:39 9.metadata
```

Empty-sized segments are continuously generated when the validator node infinitely restarts.

Impact

Low

Repeatedly generating zero-byte segments can lead to resource exhaustion by consuming inodes and disk space for corrupted segments. Although systemd's backoff mechanism aids in restart loops, it does not prevent the creation of these zero-byte segments.

Recommendation

While there is no current bug causing infinite restarts within the EigenDA's validator code itself, a defense-in-depth strategy is necessary to prevent the excessive creation of zero-byte segments if validators repeatedly restart.

Remediation

The team has a plan and strategy to resolve finding #9. The upcoming patch will also address finding #4. Their plan is outlined below.

Upon restart, the system checks if the last segment is sealed. If it's unsealed, the system iterates through the key file of this segment, verifying each key's presence in the corresponding value file.

For each key:

- If present, the key is re-inserted into the keymap and added to a "good keys" list.
- If not present, the key is added to a "bad keys" list.

If the "bad keys" list contains any entries, a new key file is created containing only the valid keys. This new file then atomically replaces the old key file. Following this, the segment is sealed, and the metadata file is updated with the correct number of keys, resolving potential metric discrepancies after a restart.

Finally, if the "good keys" list is empty after these steps, all segment files are deleted, and the old segment number becomes available for reuse.

4. Conclusion

The LittDB security audit uncovered several security and reliability vulnerabilities, ranging from informational to high severity, all requiring attention for EigenDA's robustness.

High-severity issues, specifically a weak permutation algorithm in `util.HashKey` and race conditions in shard and key file processing, present significant risks to data integrity, confidentiality, and operational consistency, demanding immediate fixes to prevent exploitation or data corruption.

Medium-severity findings, such as potential goroutine and channel leaks, could cause resource exhaustion and long-term performance issues, necessitating timely resolution to ensure system stability and reliability.

Low-severity issues involve risks from unvalidated input from potentially malicious relays. While less critical, these indicate areas needing enhanced validation and stricter security controls to prevent unauthorized manipulation.

Informational observations, including the use of `unsafe.String`, shadowed loop variables, and insufficient input validation, do not pose direct security threats but suggest areas for improving code quality, maintainability, and defensive programming practices.

In summary, immediate action on high-severity vulnerabilities is critical, followed by a systematic approach to address medium and low-severity findings. Addressing even informational items will contribute to a more robust and maintainable codebase. A prioritized remediation plan is crucial for ensuring LittDB's secure and stable operation within EigenDA.



Theori, Inc. (“We”) is acting solely for the client and is not responsible to any other party. Deliverables are valid for and should be used solely in connection with the purpose for which they were prepared as set out in our engagement agreement. You should not refer to or use our name or advice for any other purpose. The information (where appropriate) has not been verified. No representation or warranty is given as to accuracy, completeness or correctness of information in the Deliverables, any document, or any other information made available. Deliverables are for the internal use of the client and may not be used or relied upon by any person or entity other than the client. Deliverables are confidential and are not to be provided, without our authorization (preferably written), to entities or representatives of entities (including employees) that are not the client, including affiliates or representatives of affiliates of the client.

©2025. For information, contact Theori, Inc.